

Foundations of Decision Intelligence and Decision Support Systems

Jaxon Doolittle

Georgia Southwestern State University

CSCI 5310: Decision Support Systems

Prof. Alexander Yemelyanov

June 15, 2026

1. Is decision intelligence only a set of AI capabilities? Give reasons for your answer.

No, Decision Intelligence isn't just a bunch of commercial AI apps or software capabilities. It's actually a comprehensive framework built to democratize AI, integrating raw machine power with real-world decision-making theories from behavioral, social, and management sciences. While it definitely leverages AI tools like machine learning to chew through massive datasets and automate routine tasks, its core purpose is figuring out the best way to split the decision-making workload between humans and machines. This keeps humans in the loop to handle complex, creative problem-solving, making DI a massive overarching discipline rather than just a shiny new tech toolset.

2. Compare and contrast Decision Intelligence (DI) with AI. What is the relationship between these two disciplines?

You can think of Artificial Intelligence as the underlying engine, like science and tech built to make machines think and act like humans. Decision Intelligence, on the other hand, is the practical framework that actually applies those AI capabilities to improve how we make decisions in the real world. While AI handles the heavy lifting with algorithms and neural networks, DI acts as the broad umbrella that bridges those technical tools with human behavioral science. They work hand-in-hand: AI gives you the computational horsepower to extract insights from insane amounts of data, and DI provides the structure needed to turn those insights into concrete, outcome-driven actions.

3. Is it a wise idea to delegate all our decisions to technology? Describe the importance of keeping humans in the loop for decision-making.

Handing over the keys and delegating every single decision to technology is a terrible idea because current AI systems simply aren't evolved enough to handle highly complex or completely unpredictable environments. You absolutely have to keep humans in the loop. We have unique faculties that computers can't just replicate, like genuine empathy, imagination, and creative problem-solving. When architecting something like the AI tools at Shrood (my startup), letting the machine handle routine data sorting is a no-brainer. But when it comes to massive strategic pivots or navigating different ethical situations, you need human intuition. Technology is there to share the cognitive load of programmable tasks so humans can focus their energy on the critical decisions.

4. Describe the major steps of the decision-making process. Explain how technology can support each step, using relevant examples.

The decision-making process breaks down into seven main steps, and tech can step in to augment every single one of them:

- Step 1: Identifying the Problem. This is about spotting an issue that needs fixing. For example, anomaly detection algorithms can automatically flag shady network traffic, tipping off an engineer before a human even realizes there's a breach.
- Step 2: Gathering Information. Pulling data to understand the context. Web scrapers or basic decision assistance software can rip threat intelligence reports from external databases.
- Step 3: Identifying Alternatives. Figuring out possible solutions. A decision support system could quickly generate a list of patching protocols to lock down a compromised network.
- Step 4: Evaluating Alternatives. Weighing the pros and cons of each option. Decision augmentation tech can run simulative analyses to predict the downtime and effectiveness of a specific firewall configuration.
- Step 5: Choose the Best Alternative. Picking the optimal route. Machine learning recommender systems can crunch the expected utility and literally point to the best protocol to run.
- Step 6: Implementing the Decision. Actually executing the plan. Decision automation tools, like robotic process automation, can autonomously push security patches across servers without anyone lifting a finger.
- Step 7: Reviewing the Decision. Monitoring the fallout. Automated feedback loops track system stability post-patch and log the success rates to train future machine learning models.

5. Compare the satisficing and maximizing approaches to decision-making. Provide business examples of each approach and discuss how DI can help organizations balance the two.

Satisficing is all about efficiency, finding a decision that is "good enough" without burning through time and cognitive resources analyzing every piece of data. For example, when futures trading (which is something I've been working on), you might be satisfied by picking a platform like NinjaTrader just because it reliably executes orders fast enough, rather than wasting weeks evaluating every piece of software on the market.

Maximizing, on the flip side, is about hunting for the absolute best possible outcome, which requires insanely detailed analysis. If you're gearing up to sell a company for a target valuation of around \$500,000, you are going to maximize. You'll meticulously evaluate buyers like Mutare or Pindrop and run deep financial models to ensure you get the absolute highest return.

The catch is the tradeoff: satisficing sacrifices quality for speed, while maximizing kills efficiency for quality. DI helps balance the two, a process called optimizing, by using

algorithms to chew through massive datasets instantly. This gives human decision-makers the high-quality outcomes of a maximizing approach but with the speed and low effort of satisficing.

6. What forms of DI does your company or school use?

Decision Intelligence generally falls into four buckets: Decision Assistance, Decision Support, Decision Augmentation, and Decision Automation.

- **Company Application:** When I was working with NASA and MIT, the engineering pipeline heavily relied on Decision Augmentation and Decision Automation. For example, building out Retrieval-Augmented Generation to synthesize literature is straight-up decision augmentation. So, the system I built explores data and serves up specific, actionable insights. On the other hand, engineering autonomous agents like Digital Filer pushes right into decision automation, where the machine processes the environment and executes tasks with basically zero human intervention.
- **School Application:** At Georgia Southwestern State University, it's more about Decision Support. Things like degree auditing software and learning management systems gather data and lay out alternatives to help grad students and faculty make operational choices, like picking the right courses for a Master's track. So, it's a little different.

7. Distinguish between normative and descriptive theories of decision-making.

- **Normative Theory:** This theory is all about how people should make decisions. It comes from traditional economics and assumes humans are perfectly rational agents who calculate probabilities and expected utility to maximize their returns.
- **Descriptive Theory:** This looks at how people actually make decisions in the real world. Built on behavioral economics, it accepts that humans don't sit around running statistical probabilities in their heads; instead, our choices are heavily hijacked by emotions, cognitive biases, mental models, and environmental constraints.

8. Why do people often struggle to estimate probabilities accurately or make decisions according to expected utility theory?

People struggle to act purely on expected utility because our brains are biologically wired to be lazy, naturally following the "law of least effort". Firing up the brain's rational, reflective side (System 2) to calculate hard probabilities takes a lot of mental energy. To cut corners, people rely on their automatic side (System 1) and use simple heuristics like anchoring or availability to make split-second judgments. Toss in emotional influences and

cognitive biases like loss aversion, and it's no surprise that the fundamental limits of human processing power prevent us from perfectly analyzing complex data.

9. Explain the concept of noise as described by Daniel Kahneman and his colleagues. How does noise differ from bias?

Daniel Kahneman breaks down noise as the "undesirable variability in judgments of the same problem". The easiest way to tell it apart from bias is direction. Bias is a systematic error that pushes decisions in one predictable direction, like constantly underestimating project timelines because of the planning fallacy. Noise, however, is just random, unwanted scatter in the decision-making process.

- System Noise: This is the variability you get when different people judge the same case, which gets amplified in groups based on things like who speaks first or loudest.
- Occasion Noise: This happens when the same person makes different judgments on the same issue at different times, usually because of temporary stuff like mood, fatigue, or even the weather.

10. Give examples of situations when system and occasion noise affect your judgment.

- Occasion Noise: When you're running a day trading setup on Jigsaw or NinjaTrader, your judgment is incredibly vulnerable to how you're feeling that exact day. If you are staring at the charts while fatigued or frustrated, occasion noise messes with your cognitive machinery, making you hesitate on a solid trade or jump into a risky entry that you would totally avoid if you were well-rested.
- System Noise: You see this a lot when collaborating on business strategy or reviewing AI frameworks. If one person in the room voices their opinion first and does it with total confidence, it can unconsciously skew the feedback of everyone else, creating a totally unwanted variability in how the group evaluates the strategy.

11. Describe situations in which you or someone you know made decisions primarily based on intuition or emotion. Were those decisions effective?

From a biological standpoint, human decision-making is constantly being hijacked by the limbic system, which processes emotions and can literally suppress the brain's analytical networks. Because we default to System 1 to save energy, a lot of our choices are purely feeling-driven.

Take modifying a car, for instance. Deciding to put a spoiler on a car is a choice built almost entirely on emotion and aesthetic passion. If you looked at it through a strict normative economic lens, trying to maximize financial utility, it's a terrible decision. But

descriptively, it completely maximizes emotional utility and personal satisfaction, making it highly effective for what it was actually meant to do.

12. Explain the basic decision-making mechanism of computers. What physical and non-physical capabilities do computers utilize to support decision-making?

At their core, computers are built to mimic how the human brain makes decisions: they take in data, process it, and spit out an action.

- **Physical Capabilities (Hardware):** The CPU is the brain of the operation, actively running the code. RAM acts as short-term memory to handle active data, while hard drives or flash drives handle persistent, long-term storage. Bringing it all together is the motherboard, which acts like the central nervous system, using integrated circuits to shuttle data and power between all the components.
- **Non-Physical Capabilities (Software):** To actually know how to decide, computers rely on programs. This scales from basic traditional programming all the way up to deep AI models that process insanely complex datasets to execute autonomous actions.

13. What is embodied intelligence? Discuss its significance in human and machine decision-making.

Embodied intelligence is exactly what it sounds like: a cognitive system built into a physical body to interact with the world.

- **In Humans:** We are naturally embodied. We pull data in through sensors (eyes, ears) and execute our choices through actuators (hands, vocal cords).
- **In Machines:** This is where AI gets physical. Embodied intelligence powers autonomous agents, like self-driving cars or industrial robots. Instead of eyes, they use digital sensors like cameras, and instead of hands, they use mechanical actuators like motors. It's a huge deal because it allows AI to actually navigate and manipulate the messy, non-deterministic real world without needing a human to pull the levers.

14. Compare supervised and unsupervised machine learning. How do these approaches help intelligent systems make decisions?

- **Supervised Machine Learning:** This is when you train a model using labeled data, meaning you feed the system the input and the correct output. It teaches the system how to make decisions by mapping data features to specific labels, making it great for regression or classification tasks.

- Unsupervised Machine Learning: Here, you drop the system into raw, unlabeled data. It doesn't have an answer key. Instead, it uses clustering algorithms to find hidden patterns, group similar items together, or flag anomalies.
- If you're running a cybersecurity platform, supervised learning is what you use to scan for known malware signatures because the threats are clearly labeled. Unsupervised learning is what you deploy to monitor live network traffic and cluster weird behaviors to catch zero-day vulnerabilities no one has seen before.

15. How does the depth of an artificial neural network influence its ability to learn patterns and produce accurate predictions?

Artificial neural networks are designed to mimic the biological wiring of the human brain using interconnected layers of artificial neurons (perceptrons) that talk to each other via weights and activation functions. A basic, single-layer network is fine for simple linear classification, but when you stack multiple "hidden layers" into the architecture, you get deep learning. Adding that depth completely changes the game. It allows the system to process massive chunks of unstructured data, like raw audio or video. The early layers handle low-level stuff (like detecting edges), while the deeper layers synthesize those details into highly complex features. Ultimately, this depth lets the network navigate nonlinear relationships, spitting out insanely accurate predictions.

16. What is reinforcement learning? Explain how reinforcement learning systems acquire and refine behaviors through interaction with their environment.

Reinforcement Learning is a machine learning setup that lets an intelligent agent learn purely through trial, error, and practice without anyone holding its hand. Think of it like operant conditioning from psychology: the system refines its behavior by getting positive rewards for getting things right and negative penalties when it messes up. As the agent grinds through interactions with its environment, it starts to approximate the optimal value function, basically predicting the best possible moves to make in the future. If you were deploying an automated Python-based trading bot like Midas (my day trading bot), an RL algorithm would let the bot learn the futures market by rewarding it for profitable trades and penalizing it for blown accounts, gradually optimizing the strategy on its own.

17. Identify situations or environments in which humans generally make better decisions than machines. Justify your answer with examples.

Machines fall flat in environments that are unobservable, highly dynamic, or non-deterministic, basically anywhere the rules are constantly shifting or the data is hidden. Humans absolutely dominate in these scenarios because we have biological traits computers don't: imagination, creativity, and empathy. For example, if you are negotiating the sale of an AI startup, you are dealing with corporate politics, reading the room, and

guessing the unobservable motives of the buyers. An AI model can't replicate the emotional intelligence and creative deal-structuring required to get a buyout across the finish line.

18. Identify and briefly describe five technologies, systems, solutions, or applications for each of the following DI levels: decision assistance, decision support, decision augmentation, and decision automation.

- Decision Assistance: AI does the bare minimum to gather info while humans keep full control.
 1. Web-based research tools.
 2. Basic internet search engines.
 3. Third-party apps are embedded inside digital devices.
 4. GPS navigation apps.
 5. Simple web dashboards.
- Decision Support: Systems that pull data and map out alternatives for you.
 1. Customer Relationship Management platforms.
 2. Enterprise Resource Planning architecture.
 3. Clinical Decision Support Systems that recommend medical treatments.
 4. Communication-driven groupware for team collaboration.
 5. Spreadsheet-based financial models.
- Decision Augmentation: Machines get proactive, crunching data to recommend actual outcomes.
 1. Recommender systems (the algorithms powering Netflix or Spotify).
 2. AutoML platforms that train models automatically.
 3. Natural Language Processing tools for automated content tagging.
 4. AI-powered capacity command centers.
 5. Simulative analysis systems.
- Decision Automation: Machines take the wheel, making the choice and executing it.
 1. Robotic Process Automation workflows.
 2. Self-driving autonomous vehicles.
 3. Autonomous industrial robotics.

4. AI-driven robot advisors managing retail stock portfolios.
5. Automated inventory replenishment systems.

19. Identify the main features of Intelligent Decision Support Systems (IDSSs).

Intelligent Decision Support Systems are essentially high-level intelligent agents built to mimic the human decision-making process as closely as possible. Their main features include the horsepower to mine massive datasets in real-time, spot hidden trends, and autonomously generate concrete recommendations. They pull this off by wrapping advanced machine learning algorithms together with specialized techniques like case-based reasoning, fuzzy logic, and multiattribute decision-making, allowing them to produce solid results even when the environment is completely chaotic or uncertain.

20. Compare IDSSs with traditional DSSs. Discuss their similarities, differences, advantages, and limitations.

- **Similarities:** Both traditional DSSs and IDSSs are interactive systems integrated into a company's architecture to crunch data, evaluate options, and help teams make better business calls.
- **Differences & Advantages:** Traditional DSSs are pretty rigid. They rely on simple rules or manual inputs, meaning humans still have to do heavy lifting to figure out what the data actually means. IDSSs are proactive. By leveraging AI and machine learning, an IDSS can automatically adapt to complex environments, find nonlinear patterns, and spit out accurate predictions that totally eclipse the computational limits of a legacy DSS.
- **Limitations:** Traditional DSSs hit a wall because they can't process massive, unstructured datasets and still rely too heavily on human intuition, which opens the door for cognitive bias. On the other hand, IDSSs are resource-heavy, demanding insane amounts of data and compute power to train. Plus, if the engineers writing the code inject their own biases into the algorithm, the whole network becomes compromised.