

The Mechanics of Choice: Integrating Human Psychology and AI in Complex Decision Systems

Jaxon Doolittle

Georgia Southwestern State University

CSCI 5310: Decision Support Systems

Prof. Alexander Yemelyanov

June 22, 2026

Imagine trying to build a secure, enterprise-grade application without understanding how the end-user is actually going to interact with it. It would be a total mess, right? You can't just write code in a vacuum; you have to understand the human element; the flaws, the shortcuts, and the biases that dictate behavior. The exact same principle applies to decision-making. Historically, economists wanted to believe that humans were perfectly rational machines that calculated probabilities to maximize payouts. But let's be real: that's not how the world works. Humans are noisy, emotional, and heavily constrained by their own biology.

As we push further into the era of Artificial Intelligence and Decision Support Systems, understanding the gap between how we should make decisions and how we actually make them is critical. You can't just build a system and ignore what happens when a stressed-out user gets their hands on it. Everything has to be done with purpose and accountability. This report breaks down the mechanics of human choice. We will look at everything, from the foundational models of decision theory to the cognitive biases that derail us, and finally, to how AI and decision support tech can offset our broken wetware to solve real-world problems.

1. The Integrative Model of Decision-Making: Navigating Conflict

The traditional view of decision-making often treats it as a simple input-output mechanism. But in reality, decision-making is a complex stage of information processing that sits right between perceiving a stimulus and executing a response (Lehto & Nanda). To map out this reality, Lehto and Nanda present the Integrative Model of Decision-Making.

In this model, human decisions are heavily dictated by the decision context (e.g., time pressure, single vs. multiple decision-makers) and the presence of conflict. In a normal, highly routine situation, a decision-maker perceives a stimulus, recognizes the pattern, and executes an action instantly. Think of a seasoned developer seeing a syntax error and automatically hitting the backspace key. That is an associative leap. But when the environment becomes uncertain, conflict is introduced, acting as a hard roadblock that forces the brain out of autopilot.


Conflict is basically the trigger for non-routine decision-making. It arises from several layers of uncertainty:

- Unidentified conditions: You don't even know what the problem is or what your alternatives are.
- Uncertain preferences/consequences: You don't know what will happen if you pull the trigger on an action, or you don't even know what outcome you actually prefer.
- Conflicting objectives: You want a software patch to be 100% secure, but you also need it deployed in under an hour. You can't have both.
- Lack of consensus: When dealing with groups, different stakeholders want different things.

To move past this conflict, the decision-maker must use judgment to resolve the uncertainty. If consequences are uncertain, they must judge the likelihoods (probabilities). If objectives conflict, they must judge importance and assign weights. If there is a lack of consensus, they have to debate or vote.

Once the conflict is resolved, the decision-maker frames the problem and selects a decision rule (like satisficing or maximizing). The rule tells them how to evaluate the alternatives and finally execute the action.

Figure 1: Simplified Flow of the Integrative Model of Decision-Making



2. The Battle of Decision Approaches: From Perfection to Pragmatism

Once a decision is made, the user has to apply a decision rule. The problem is that not all rules are created equal. Some demand insane amounts of cognitive energy, while others are quick and dirty. Lehto and Nanda outline several major approaches, which drastically differ in how they treat rationality, preferences, and uncertainty.

Subjective Expected Utility Theory:

SEU is the "perfect world" normative model. It assumes the decision-maker is a completely rational agent who can perfectly quantify their preferences into a utility function. Under SEU, a user looks at an uncertain event, assigns a subjective probability to every possible consequence, multiplies the utility by that probability, and chooses the option that yields the absolute highest expected return (Lehto & Nanda). It treats uncertainty as something that can be neatly calculated using Bayesian math. The problem? It requires heavy lifting that the human brain just isn't built to do consistently.

Multi-attribute Utility Theory:

MAUT takes SEU and scales it up for the real world, where decisions rarely have just one objective. Instead of just looking at financial payout, MAUT allows a decision-maker to evaluate multiple conflicting objectives (e.g., cost vs. performance vs. security). It breaks down the decision into a value tree, assigning specific weights to different attributes so they can be mathematically compared (Lehto & Nanda). It's a compensatory model, meaning a high score in one area (like

ultra-high security) can compensate for a low score in another (like a terrible user interface).

Dominance:

Dominance is the most fundamental, no-brainer normative rule. It occurs when one alternative is objectively better (or at least equal) to another alternative across every single possible event or attribute (Lehto & Nanda). If Software A is cheaper, faster, and more secure than Software B, Software A dominates. The treatment of uncertainty here is irrelevant; because the choice is better in every scenario, no complex probability math is required.

Satisficing:

If SEU and MAUT are about maximizing, satisficing is about efficiency. Developed by Herbert Simon, satisficing acknowledges that people don't have the time or brainpower to evaluate every single option on the market. Instead of looking for the "best" outcome, the decision-maker sets a minimum aspiration level (a baseline threshold). They evaluate options sequentially and pull the trigger on the very first alternative that meets those minimum requirements (Lehto & Nanda). It treats rationality pragmatically: it is a non-compensatory strategy that stops the search early to save cognitive resources.

Table 1: Comparison of Decision-Making Approaches

Approach	Rationality Assumption	Treatment of Preferences	Treatment of Uncertainty	Multiple Objectives
SEU	Purely rational (Maximizer).	Preferences translated directly to a numerical utility curve.	Handled via subjective probability & Bayesian math.	Generally single-objective focused.
MAUT	Purely rational (Maximizer).	Preferences are weighted based on importance.	Handled via probability, applied across a value tree.	Explicitly designed to trade-off conflicting objectives.
Dominance	Highly rational, but conceptually simple.	Strict preference; one option must be universally better.	Bypassed; the option wins regardless of the environment's state.	Requires the option to win/tie across all objectives.

Satisficing	Boundedly rational (Efficiency-focused).	Preferences simplified to a "minimum acceptable" baseline.	Ignored/minimized; search stops before full environment is mapped.	Non-compensatory; must meet the baseline for all required objectives to be chosen.
-------------	--	--	--	--

3. Bounded Rationality, Heuristics, and Our Broken Brains

Classical economic models like SEU assume we are all walking supercomputers. But behavioral economics proves that actual human decision-making is broken from the ground up when compared to these rigid normative assumptions.

The core of this is Bounded Rationality, a concept by Herbert Simon which states that human rationality is strictly limited by the information we have, the time we are given, and the processing capacity of our brains (Lehto & Nanda). Because our brains follow the "law of least effort," we physically cannot process massive, multi-variable decision trees in our heads. To survive, we use cognitive shortcuts or rules of thumb that let us make split-second judgments.

Tversky and Kahneman identified several major heuristics that humans use to bypass complex math:

- **Representativeness:** Judging the probability of something based on how closely it resembles a familiar prototype. For instance, an employee might assume an email is safe because it looks representative of a corporate IT memo, completely ignoring the base-rate probability of it being a phishing attack.
- **Availability:** Judging the likelihood of an event based on how easily it comes to mind. After a massive, highly publicized ransomware attack like Colonial Pipeline, executives will wildly overestimate the immediate probability of their own company being targeted because the memory is highly salient.
- **Anchoring:** Starting with an initial piece of information (the anchor) and failing to adjust adequately. If a software vendor quotes an enterprise license at \$100,000, any counteroffer around \$80,000 feels like a win, even if the software is objectively only worth \$40,000.

While heuristics are ecologically rational and save us time, they create cognitive biases, systematic predictable errors in judgment. Because we rely on shortcuts, we fall victim to the Confirmation Bias (selectively seeking out data that proves we are right while ignoring red flags), the Status Quo Bias (refusing to upgrade legacy tech because changing is too mentally taxing), and the Dunning-Kruger Effect (overestimating our own technical competence) (Lehto & Nanda). Actual human decision-making is messy. It doesn't follow Bayesian formulas; it follows the path of least resistance.

4. Context Changes Everything: Behavioral, Naturalistic, and Group Dynamics

The environment in which a decision is made completely alters how the brain operates. Lehto and Nanda highlight three distinct approaches to understanding contextual decision-making: Behavioral, Naturalistic, and Group.

Behavioral Decision-Making:

This approach bridges psychology and economics, focusing on how framing and emotion dictate choice. In a strictly rational model, a loss of \$1,000 and a gain of \$1,000 have the same absolute value. But behaviorally, humans suffer from loss aversion, the psychological pain of losing is significantly stronger than the joy of gaining. Furthermore, Tversky and Kahneman's Prospect Theory proved that decisions are completely manipulated by how they are framed. If a surgeon frames a procedure as having a "90% survival rate," a patient will likely agree. If framed as having a "10% mortality rate," they will likely refuse, even though the statistical reality is exactly the same (Lehto & Nanda).

Naturalistic Decision-Making (NDM):

NDM looks at how actual experts, like firefighters, ER doctors, or incident response engineers make decisions in the trenches under extreme time pressure and stress. In these environments, you do not have time to sit down and build a MAUT value tree. Gary Klein's Recognition-Primed Decision model shows that experts rarely compare alternatives. Instead, they use their massive reservoir of experience to perceive cues, recognize a pattern, and immediately execute a rule-based response (Lehto & Nanda). They simulate the action in their head; if it works, they do it. Expertise allows for rapid, associative leaps, bypassing the heavy cognitive load of analytical decision-making entirely.

Group Decision-Making:

When you put multiple people in a room, you introduce a whole new layer of chaos: social norms, ethics, and power dynamics. Conflict in groups can be either substantive (intellectual clashes over the actual data, which improves decision quality) or affective (emotional, ego-driven clashes, which destroy it) (Lehto & Nanda). Groups are also highly susceptible to Groupthink and Herd Behavior, where the desire for social conformity causes the group to blindly agree with the highest-status individual or the loudest voice, completely ignoring critical red flags. While groups have the potential for higher "collective intelligence," the social pressure to conform often leads to riskier, more polarized decisions than an individual would make on their own.

5. The Role of AI-Assisted Decision Support as the Great Equalizer

If humans are inherently biased, easily fatigued, and incapable of processing massive datasets, we need a bridge to offset those limitations. That is exactly where Decision Support Systems and Artificial Intelligence step in. The role of AI-assisted DSS is to take the heavy lifting of

normative models (like SEU and MAUT) and automate them so humans can actually reap the benefits without burning out their cognitive resources.

A traditional DSS helps organize data, but an Intelligent Decision Support System acts as a proactive advisor. Using machine learning models, like Support Vector Machines, Naïve Bayes, and Deep Neural Networks, an IDSS can ingest terabytes of unstructured data, find nonlinear patterns, and classify risks with a level of accuracy that humans literally cannot achieve (Lehto & Nanda).

For example, in a modern enterprise, Information Overload is a critical threat. Presenting an analyst with millions of rows of log data leads to decision fatigue and choice deferral. AI fixes this through Visual Analytics. Tools like Microsoft Azure or Tableau synthesize ambiguous, massive datasets into interactive visual interfaces (like heat maps or node graphs). Because vision is our highest-bandwidth sensory channel, visual analytics allows a human operator to detect anomalies instantly. Furthermore, AI helps strip out occasion noise (fatigue, bad moods) and confirmation bias by forcing the human operator to look at statistically weighted, objective recommendations rather than relying on their gut.

6. Real-World Application: Architecting the Midas Trading Bot

To see how these theories actually collide in the real world, we can look at the architecture of "Midas," an automated Python-based algorithmic day-trading bot I developed to trade Micro E-mini futures. Financial markets are the ultimate non-deterministic, highly dynamic environments. The rules constantly shift, information is partially unobservable, and the risk of financial ruin is omnipresent.

The Human Problem (Applying Satisficing):

When a human day-trader sits at a terminal, they are bombarded with information overload: Level 2 order flow, VWAP (Volume Weighted Average Price) bands, RSI indicators, and live macroeconomic news. According to Simon's bound rationality, the human brain cannot process all of this simultaneously. Furthermore, human traders are severely impacted by occasion noise, if a trader is fatigued or frustrated from a previous loss, their cognitive machinery breaks down.

As a result, a human trader operating under extreme time pressure rarely maximizes. Instead, they apply the Satisficing decision rule. They set a minimum aspiration level (e.g., "If the RSI is below 30 and price touches VWAP, I will buy"). They stop scanning the broader market context and pull the trigger on the first setup that looks "good enough." Because of the Availability Heuristic, they might also heavily overweight recent price action while ignoring larger macroeconomic trends, leading to suboptimal entries and blown accounts.

The Algorithmic Solution (Applying MAUT):

To fix this, the Midas bot was engineered to strip out the human element entirely. Instead of satisficing, Midas applies a strict Multi-attribute Utility Theory (MAUT) framework. Before executing a trade, the bot pulls in data across multiple conflicting attributes: order book imbalance, tick

volatility, trailing momentum, and dynamic risk-to-reward ratios. Midas applies specific weights to each of these attributes, calculates the expected utility of going long versus going short, and executes the trade only if the additive utility crosses a strict mathematical threshold.

How AI-Assisted Decision Support Improves the Process:

Building an IDSS architecture around Midas completely revolutionizes the decision process. By utilizing Reinforcement Learning, Midas acts as an autonomous intelligent agent.

1. **Elimination of Cognitive Bias:** The AI does not suffer from loss aversion or the gambler's fallacy. If a trade goes against the bot, it doesn't "revenge trade" out of anger; it simply recalculates the expected utility based on the new market state and cuts the loss precisely where the algorithm dictates.
2. **Overcoming Bounded Rationality:** While a human trader can realistically monitor maybe two or three asset classes at once, the AI-assisted system can monitor thousands of data points across the S&P 500, Nasdaq, and Russell 2000 indices simultaneously, scanning for dominance and arbitrage opportunities in milliseconds.
3. **Dynamic Adaptation:** Through reinforcement learning, the bot constantly refines its own MAUT weights based on environmental feedback. It learns which attributes yield the highest profitability in different market regimes without requiring a human to manually rewrite the code.

By integrating AI into the decision-making loop, we successfully bridge the gap between descriptive reality and normative perfection. We take a chaotic, noise-filled human process and replace it with a system built on purpose, accountability, and ruthless mathematical precision.

References

Lehto, M. R., & Nanda, G. (2020). Decision-Making Models, Decision Support, and Problem Solving. In G. Salvendy & W. Karwowski (Eds.), *Handbook of Human Factors and Ergonomics* (5th ed., pp. 159-202). Wiley.