

Navigating Complex Choices with AHP and MAUT

Jaxon Doolittle

Georgia Southwestern State University

CSCI 5310: Decision Support Systems

Assignment: Week 3 Submission

Prof. Alexander Yemelyanov

Due: July 6, 2026

Submitted: June 30, 2026

When dealing with high-stakes decisions, whether it's provisioning enterprise servers, buying a car, or designing an algorithmic trading bot, relying on a "gut feeling" is a recipe for disaster. The human brain simply cannot hold dozens of conflicting variables in its working memory. Multi-Criteria Decision Analysis frameworks like the Analytic Hierarchy Process and Multi-Attribute Utility Theory act as external processors for our brains. They force us to break down massive, chaotic decisions into structured, mathematical models. This report unpacks the theory behind these models and applies them to real-world scenarios using Super Decisions and RightChoice software.

1. Analytic Hierarchy Process and Super Decisions

1.1 AHP Theory & Concepts

a) The Role of Pairwise Comparisons

Trying to rank ten different alternatives simultaneously based on five different criteria requires an insane amount of cognitive load. Humans are terrible at absolute judgments, but we are highly effective at relative judgments. Pairwise comparisons solve this by utilizing a "divide and conquer" philosophy. By forcing the decision-maker to compare only two elements at a time (e.g., "Is Price more important than Maintenance, and by how much on a 1-9 scale?"), AHP drastically reduces cognitive fatigue and yields much more reliable, objective data than trying to eyeball the entire board at once.

b) Eigenvalues and Eigenvectors in AHP

Once you input your pairwise comparisons into a matrix, AHP has to extract the actual priority weights from your subjective ratings. It does this by calculating the principal eigenvalue and its corresponding eigenvector from the comparison matrix. In practical terms, the eigenvector acts as the mathematical "center of gravity" for your judgments. It synthesizes all the individual A-vs-B comparisons into a single, normalized priority vector that dictates the overall weight of each criterion.

c) The Inconsistency of Human Judgment

Humans are naturally inconsistent. We are notorious for intransitive logic. For example, a user might say that feature A is better than B, and B is better than C, but then turn around and say C is better than A. Because pairwise comparisons can easily trap a user in these contradictory loops, AHP uses a Consistency Ratio. The CR measures how far the user's judgments deviate from pure logical consistency. A CR of less than 0.1 (10%) is generally considered acceptable, acting as a built-in guardrail to ensure the underlying data isn't garbage.

d) Sensitivity Analysis

A decision isn't truly solid if breathing on the variables causes the whole model to collapse. Sensitivity analysis is essentially a stress test for the decision's robustness. It allows the decision-maker to dynamically tweak the weights of the criteria (e.g., shifting the importance of "Price" up by


5%) to see if the final ranking of alternatives changes. If a minor shift in criteria weights completely flips the winner, the decision is highly volatile.

1.2 Applying AHP Using Super Decisions (Automobile Selection)

To test AHP in a practical environment, I modeled a car-selection decision comparing three vehicles: a Truck, a SUV, and a Sedan.

The Hierarchy (4-Level Structure):



- Level 1 (Goal): Select the Best Car.
- Level 2 (Criteria): Price, Safety, Comfort, and Performance.
- Level 3 (Sub-Criteria): Purchase Cost, Operating Cost, Crash Rating, and Handling.
- Level 4 (Sub-Sub-Criteria): Down Payment, Monthly Payment, Fuel, and Maintenance.
- Level 5 (Alternatives): Sedan, SUV, Truck.



Information Panel	Network	Judgments	Results
Info 0 Node Cluster Attachments Model Structure Create/Edit Details Show Priorities Make/Show Connections	1. Choose Choose Node Select Best Car Cluster: Goal Choose Cluster Criteria	2. Node comparisons with respect to Select Best Car Graphical Verbal Matrix Questionnaire Direct Comparisons wrt "Select Best Car" node in "Criteria" cluster Safety is equally to moderately more important than Price 1. Comfort >=9.5 9 8 7 6 5 4 3 2 2 3 4 5 6 7 8 9 >=9.5 No com 2. Comfort >=9.5 9 8 7 6 5 4 3 2 2 3 4 5 6 7 8 9 >=9.5 No com 3. Comfort >=9.5 9 8 7 6 5 4 3 2 2 3 4 5 6 7 8 9 >=9.5 No com 4. Performance >=9.5 9 8 7 6 5 4 3 2 2 3 4 5 6 7 8 9 >=9.5 No com 5. Performance >=9.5 9 8 7 6 5 4 3 2 2 3 4 5 6 7 8 9 >=9.5 No com 6. Price >=9.5 9 8 7 6 5 4 3 2 2 3 4 5 6 7 8 9 >=9.5 No com	3. Results Inconsistency: 0.08624 Hybrid Comfort 0.13922 Performance 0.51194 Price 0.08665 Safety 0.26219

Consistency Ratios and Final Rankings

Throughout the matrix comparisons, the software continuously tracked the Consistency Ratios. By ensuring careful, logical comparisons, the major matrices maintained a CR well below the 0.1 threshold, confirming the validity of the inputs.



Discussion and Recommendation

Based on the synthesized priorities, the Sedan emerged as the recommended automobile. While the other vehicles outranked in each other in sub-criteria, the global weights placed a massive premium on the Performance and safety criteria. Because the Sedan scored exceptionally high in performance and safety, it mathematically eclipsed the benefits of the alternatives.

Sensitivity Analysis

To stress-test this decision, I ran a sensitivity analysis on the Performance criterion. By artificially inflating the weight of Comfort, the Sedan still won. However, if another option was selected, a different vehicle could have easily overtaken the Sedan. This proves that while the current model is robust under performance-focused weights, a slight shift in user preference completely alters the optimal outcome.

2. Multi-Attribute Utility Theory (MAUT) and RightChoice

2.1 MAUT Theory & Concepts

a) The Purpose of the Utility Function

In the real world, you can't directly compare raw horsepower to a dollar amount, the units are totally incompatible. The utility function acts as a universal translator. It converts raw, abstract performance metrics into a standardized "satisfaction meter" (a score between 0 and 1). This allows a decision-maker to accurately measure how much personal value they actually extract from a specific feature.

b) Marginal Utility Functions and Criterion Weights

Marginal utility functions map exactly how satisfaction scales with performance, capturing the reality of diminishing returns. For example, upgrading from a 20-inch TV to a 50-inch TV provides massive utility, but upgrading from an 80-inch to an 85-inch provides almost none. Exponential utility functions capture this curve perfectly. Criterion weights, on the other hand,

force the decision-maker to prioritize. They act as multipliers, ensuring that a high utility score in a useless category doesn't artificially inflate an alternative's overall rank.

c) Ranking Alternatives with Utility Scores


MAUT uses a descriptive additive model. Once the software calculates the marginal utility for every individual criterion, it multiplies each score by its assigned criterion weight and adds them all together. This mathematical aggregation squashes multi-dimensional criteria into a single, bulletproof global utility score, making it incredibly easy to rank alternatives objectively from highest to lowest.


d) Sensitivity and "What-If" Analysis

In decision support, a static model is dangerous because the real world is highly dynamic. "What-if" analysis is the ultimate risk management tool. It allows decision-makers to test the boundaries of their logic by asking, "What if the budget gets slashed?" or "What if customer reviews become twice as important?" It exposes the tipping points in a decision, highlighting exactly where an alternative goes from being optimal to obsolete.

2.2 Applying MAUT Using RightChoice

To demonstrate MAUT in action, I solved the smartphone-selection problem (Case Study 4.1) evaluating alternatives SP1 through SP5 across four criteria: Price, Customer Review, Screen Size, and Storage Size.





Marginal and Global Utility Scores (Task e)


By forcing the Price and Customer Review data through an exponential function with an exponent of 20, the utility curves became highly aggressive. The scores dropped off rapidly for sub-optimal inputs compared to the linear scales used for Screen and Storage Size.

Discussion on Utility Functions and Criterion Weights

The choice to use an exponential curve rather than a linear one drastically affected the final results. Because the Price criterion operated on a negative exponential curve, smartphones that were only slightly more expensive than the baseline suffered a massive penalty in their marginal utility score. When combined with the heavy weights distributed to Price (0.35) and Reviews (0.35), alternatives like SP2 (which had a high price but massive storage) were brutally mathematically punished, dropping straight to the bottom of the rankings.

Sensitivity Analysis Findings (Tasks f & g)

The sensitivity analysis on the Price criterion demonstrates that the decision is fairly robust. SP3 remains the top-ranked smartphone across a wide range of weights. However, if the weight placed on the Price criterion increases to approximately 0.82, a rank reversal occurs, and SP4 overtakes SP3 to become the preferred alternative.



Score Analysis Benchmark (Task h)

Based on the Score Analysis, SP2 did not reach the benchmark of 20 for the Price criterion.

